

# Regularization and Convolutional Neural Networks

---

**Deep Learning course - SKEMA 2025**

**Mastère Spécialisé® Chef de Projet Intelligence Artificielle**

**Salem Lahlou**

Some figures adapted from 3Blue1Brown (YouTube)

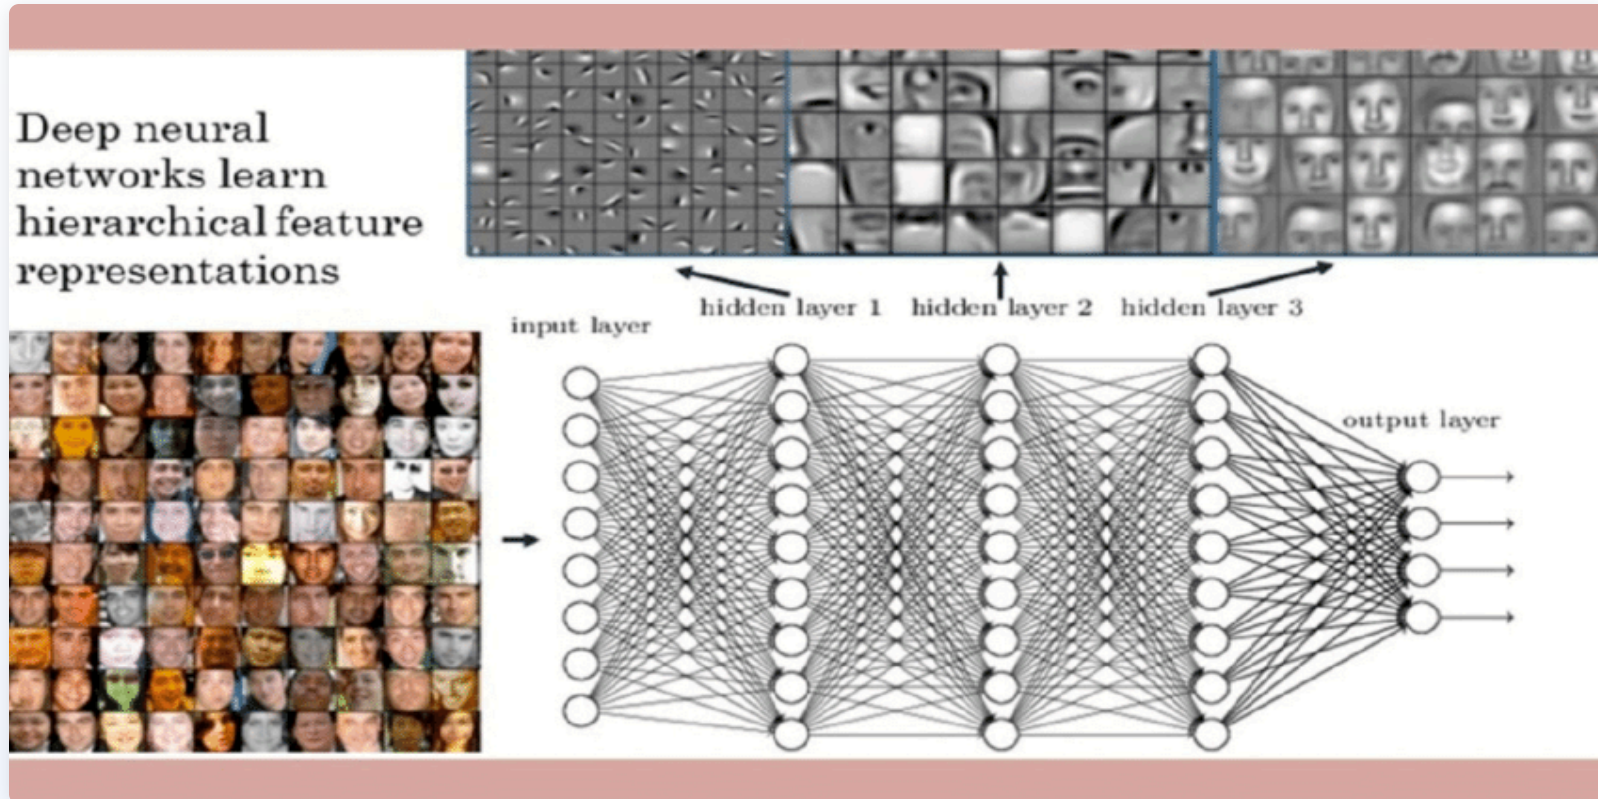
## **Shallow Neural Networks:**

- One or few hidden layers
- Limited representation power

## **Deep Neural Networks:**

- Many hidden layers
- Increased capacity to learn complex patterns

**Key Insight:** Each layer learns increasingly abstract features



**Example in image recognition:**

- First layers: Edges and simple shapes
- Middle layers: Parts and textures
- Deep layers: Complex objects and concepts

## What is overfitting?

- Model performs well on training data
- Model performs poorly on new, unseen data
- Model learns noise instead of true patterns

## Signs of overfitting:

- Large gap between training and validation error
- Model complexity increases but validation performance worsens
- Perfect or near-perfect training accuracy

## Why it happens:

- Too many parameters relative to data points
- Training too long
- Insufficient regularization
- Noisy training data

## 1. Data-based approaches:

- More training data
- Data augmentation
- Noise injection
- Feature selection

## 2. Model complexity reduction:

- Simpler models (fewer layers/neurons)
- Early stopping
- Pruning

## 3. Regularization techniques:


- L1 regularization (Lasso):  $\lambda \sum_{i=1}^p |\theta_i|$
- L2 regularization (Ridge):  $\lambda \sum_{i=1}^p \theta_i^2$
- Dropout: Randomly "drop" neurons during training
- Batch normalization: Normalize layer inputs

## Concept:

- Monitor validation error during training
- Stop when validation error starts to increase
- Take model parameters from best validation point

## Implementation:

- Track validation metrics after each epoch
- Save model whenever validation improves
- Use patience parameter to allow for fluctuations



Source: Analytics Vidhya

## **Specialized neural networks for specific data types:**

- Convolutional Neural Networks (CNNs) for images and spatial data
- Recurrent Neural Networks (RNNs) for sequential data
- Transformers for sequential data with attention mechanisms

## **Today we'll focus on CNNs:**

- Inspired by visual processing in the brain
- Excellent for image recognition and computer vision
- Leverages spatial structure in data

## Why are images difficult for standard neural networks?

- High dimensionality (millions of pixels)
- Spatial relationships matter
- Same object can appear in different positions
- Lighting, angle, and scale variations

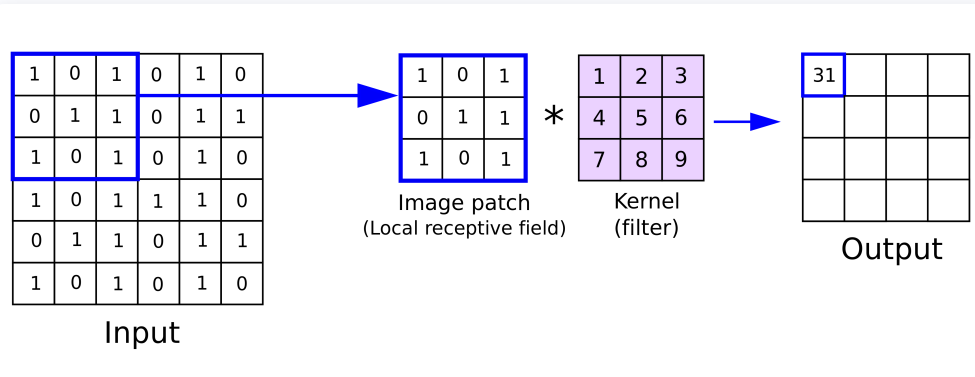
## Example:

- A 256x256 RGB image has 196,608 pixels ( $256 \times 256 \times 3$ )
- A fully connected first layer would need millions of parameters



# CNNs: Inspired by the Visual Cortex

- Neurons in visual cortex respond to specific regions of visual field
- Different neurons detect different features (edges, shapes, etc.)
- CNNs mimic this structure with specialized layers



Source: Anh Reynolds

## Advantages:

- Drastically reduces parameters compared to fully connected
- Preserves spatial relationships
- Translation invariance

## Key innovation: Parameter sharing and local connectivity

### How convolution works:

- Filters (kernels) slide across the input image
- Each filter performs the same operation at each position
- Captures local patterns regardless of position

### Mathematical representation:

- For a filter  $W$  applied to image region  $X$ :
- Output =  $\sum_{i,j} W_{i,j} \cdot X_{i,j} + b$
- Creates a feature map highlighting where patterns appear


### Example: Edge detection

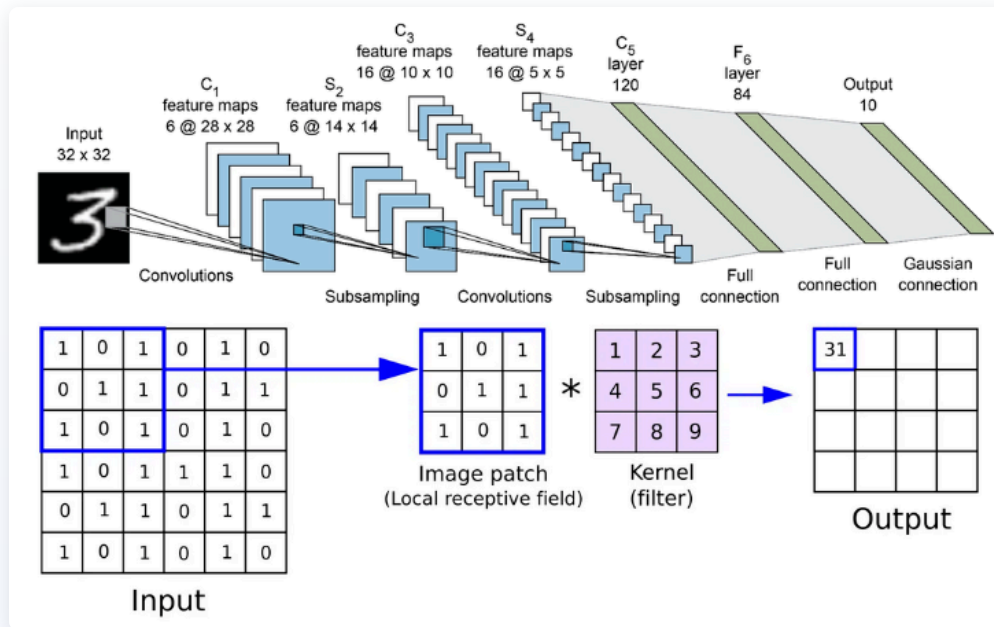
- Filter:  $\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$
- Detects vertical edges
- Different filters detect different patterns

### Goal of learning (with backpropagation): Finding the best filters

## Pooling Layer:

- Reduces the size of feature maps
- Provides some position invariance
- Reduces computation and prevents overfitting
- Common approach: Max pooling (take maximum value in region)





## Typical CNN architecture:

1. Input layer (image)
2. Convolutional layer + activation
3. Pooling layer
4. Repeat 2-3 several times
5. Flatten layer
6. Fully connected layers
7. Output layer

## Each convolutional layer:

- Uses multiple filters to detect different patterns
- Creates multiple feature maps

## LeNet-5 (one of the first CNNs):

- Developed by Yann LeCun in 1998
- Used for digit recognition (MNIST dataset)
- Structure:
  - Input:  $32 \times 32$  grayscale image
  - Conv1: 6 filters of size  $5 \times 5$
  - Pool1:  $2 \times 2$  max pooling
  - Conv2: 16 filters of size  $5 \times 5$
  - Pool2:  $2 \times 2$  max pooling
  - FC1: 120 neurons
  - FC2: 84 neurons
  - Output: 10 neurons (one per digit)


## Deep CNNs that revolutionized computer vision:

### AlexNet (2012):

- Won ImageNet competition, sparked deep learning revolution
- 8 layers, 60 million parameters
- Used ReLU activation, dropout regularization

### ResNet (2015):

- Introduced residual connections to solve vanishing gradient
- Enabled much deeper networks (up to 152 layers)
- Structure:  $F(x) + x$  instead of just  $F(x)$



# Jupyter notebook time!

---

Go to the course homepage: <https://la7.lu/sk25>

Click on "Colab 2"

## Concept:

- Take a pre-trained CNN (on massive dataset like ImageNet)
- Remove the final classification layer
- Add new layers for your specific task
- Fine-tune on your smaller dataset

## Benefits:

- Requires much less data
- Much faster training
- Better performance

## Common approach:

- Freeze early layers (generic features)
- Only train later layers (task-specific features)



## **Iterative process:**

1. Start simple (baseline model)
2. Analyze errors
3. Incrementally add complexity
4. Monitor validation performance
5. Apply regularization as needed
6. Repeat until satisfactory performance

## **Hyperparameter tuning:**

- Learning rate
- Network architecture (layers, neurons)
- Regularization strength
- Optimization algorithm

## Common issues and solutions:

### Model doesn't learn (flat loss):

- Check for errors in data preprocessing
- Verify loss function implementation
- Try a larger learning rate

### Model learns then plateaus:

- Try a different architecture
- Add regularization
- Implement learning rate schedules

### Model is unstable (loss jumps around):

- Decrease learning rate
- Try gradient clipping
- Use a different optimizer (Adam)

1. Deep networks learn hierarchical representations of data
2. Convolutional Neural Networks excel at image-related tasks
3. Modern architectures solve training challenges in deep networks
4. Transfer learning enables effective use with limited data

## Quiz time

---

Go to the course homepage: <https://la7.lu/sk25>

Click on "Quiz 5"